

✔ Congratulations! You passed!

Grade received 100% To pass 80% or higher

[Go to next item](#)

Large Scale Machine Learning

Latest Submission Grade 100%

1. Suppose you are training a logistic regression classifier using stochastic gradient descent. You find that the cost (say, $\text{cost}(\theta, (x^{(i)}, y^{(i)}))$, averaged over the last 500 examples), plotted as a function of the number of iterations, is slowly increasing over time. Which of the following changes are likely to help?

1 / 1 point

✔ Correct

2. Which of the following statements about stochastic gradient descent are true? Check all that apply.

1 / 1 point

✔ Correct

3. Which of the following statements about online learning are true? Check all that apply.

1 / 1 point

✔ Correct

4. Assuming that you have a very large training set, which of the following algorithms do you think can be parallelized using map-reduce and splitting the training set across different machines? Check all that apply.

1 / 1 point

✔ Correct

5. Which of the following statements about map-reduce are true? Check all that apply.

1 / 1 point

✔ Correct